

ESE 680-004

Learning and Control

Instructor: Nikolai Matni

Class: Tu/Th 3-4.30pm in Moore 212

Office hours: Levine 374, by appointment

Website: <https://nikolaimatni.github.io/courses/ese680-fall2019/>

E-mail: nmatni@seas.upenn.edu (start subject line with ESE680)

Course overview

- We will provide a broad overview of the emerging area at the intersection of machine learning and control, with a focus on *continuous* control problems.
- **Tentative topics list:** system identification, statistical learning theory, model based control of learned systems, model free methods, safe learning and control.
- **Official prereqs:** ESE 500 (Linear Systems) and one of CIS 520 (Machine Learning) or ESE 605 (Convex Optimization).
- **Unofficial prereqs:** mathematical maturity in linear algebra, functional analysis, probability theory, and optimization. Exposure to control and machine learning a plus.
- **Main focus:** algorithms/results with *strong theoretical guarantees* about stability, performance, robustness, and sample efficiency.

Background poll

ESE 680-004 Background Poll

Please fill out background poll at
<https://www.surveymonkey.com/r/XFSR26N>

If you are sitting in, please still fill this out.

Also please send me an e-mail so I can add
you to class mailing list.

1. Optimal Control



2. Robust Control



3. Convex Optimization



4. Statistical Learning Theory



5. Reinforcement Learning



Done

Powered by
 SurveyMonkey
See how easy it is to [create a survey](#).

Logistics and grading

- **Class structure:** Roughly 50% lectures, 50% student led discussions
- **Grading:** this is a graduate level topics class, so if you show up prepared, participate, and take the assignments seriously, you will get an A.
- **Assignments:** roughly, each of you will have to scribe one lecture, and/or present one paper in class (this may change depending on numbers)
- **Project:** done individually or in groups of two. Students are encouraged to propose a topic that connects class material to aspects of their research.
 - Proposal (1-2 pages), Midterm update (2-3 pages), Final presentation, Final report (6-8 pages).
 - Meant to mimic a conference submission: intro, lit review, problem formulation, main results, experiments, discussion/conclusion

Scribing

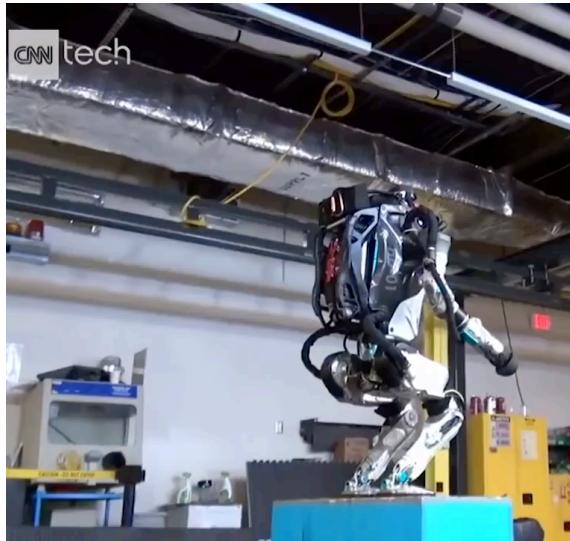
- Google doc signup sheet: https://docs.google.com/spreadsheets/d/1Kqa-ZnngOr7-c0rTKNSaCNiz_hTQARH9V_Ug8blp3Q/edit?usp=sharing
- Produce a high quality set of notes
- Go beyond what was presented in class: summarize both the lecture and the required reading. Document should be self contained.
- Use LaTeX template from course website
- **If you are sitting in, you may be asked to scribe, depending on numbers**

Paper presentations

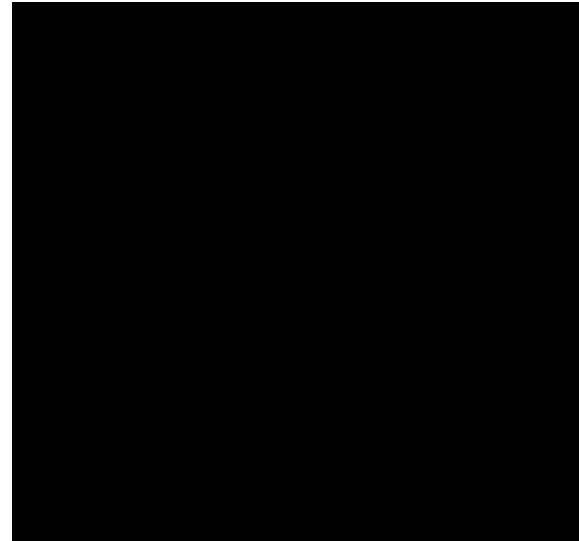
- Google doc signup sheet: https://docs.google.com/spreadsheets/d/1Kqa-ZnnqO-r7-c0rTKNSaCNiz_hTQARH9V_Ug8blp3Q/edit?usp=sharing
- Sign up to present a paper at least 2 weeks in advance, and then e-mail to schedule a meeting with me 1 week in advance to go over your presentation. Plan for about 45min of presenting and 45min of discussion.
- **Problem statement:** what problem is the paper solving, and why is it meaningful?
- **Prior work:** how does this paper fit into the current research landscape.
- **Key idea:** what is the main takeaway from the paper?
- **Key technical tools:** what are the main technical contributions of the paper? Be prepared to present and explain these in detail.
- **Points of confusion:** are there technical or conceptual arguments that are unclear or incomplete? Bring these up so that we can try to work through them in class.
- **Shortcomings/areas for improvement:** what are some of the flaws of the paper? What are possible directions for future work to address these?
- **If you are sitting in, you may be asked to present, depending on numbers**

Motivation and hand-wavy overview

The future is now



Boston Dynamics



Open AI

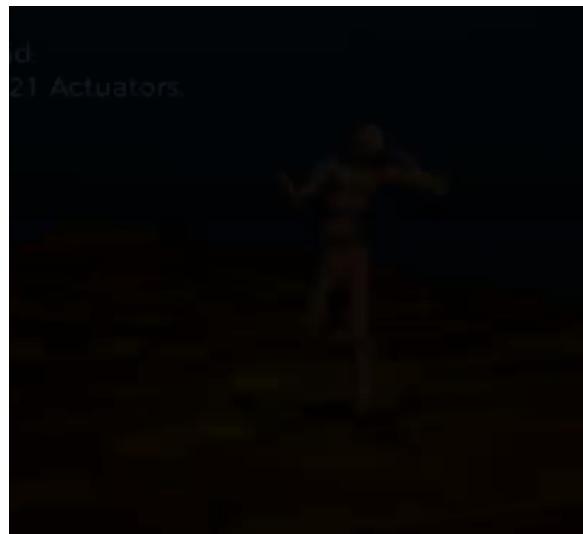


Waymo

The future is now -ish



Boston Dynamics

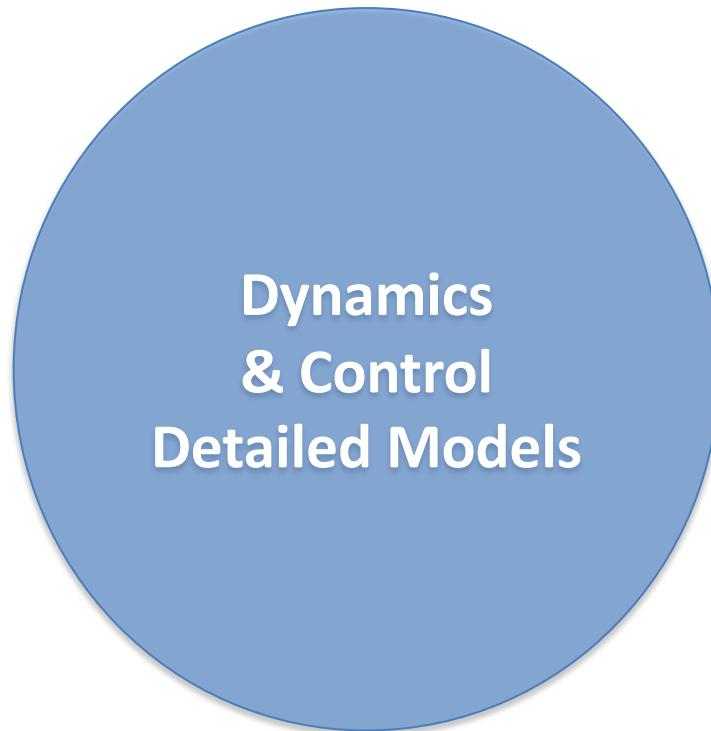


DeepMind

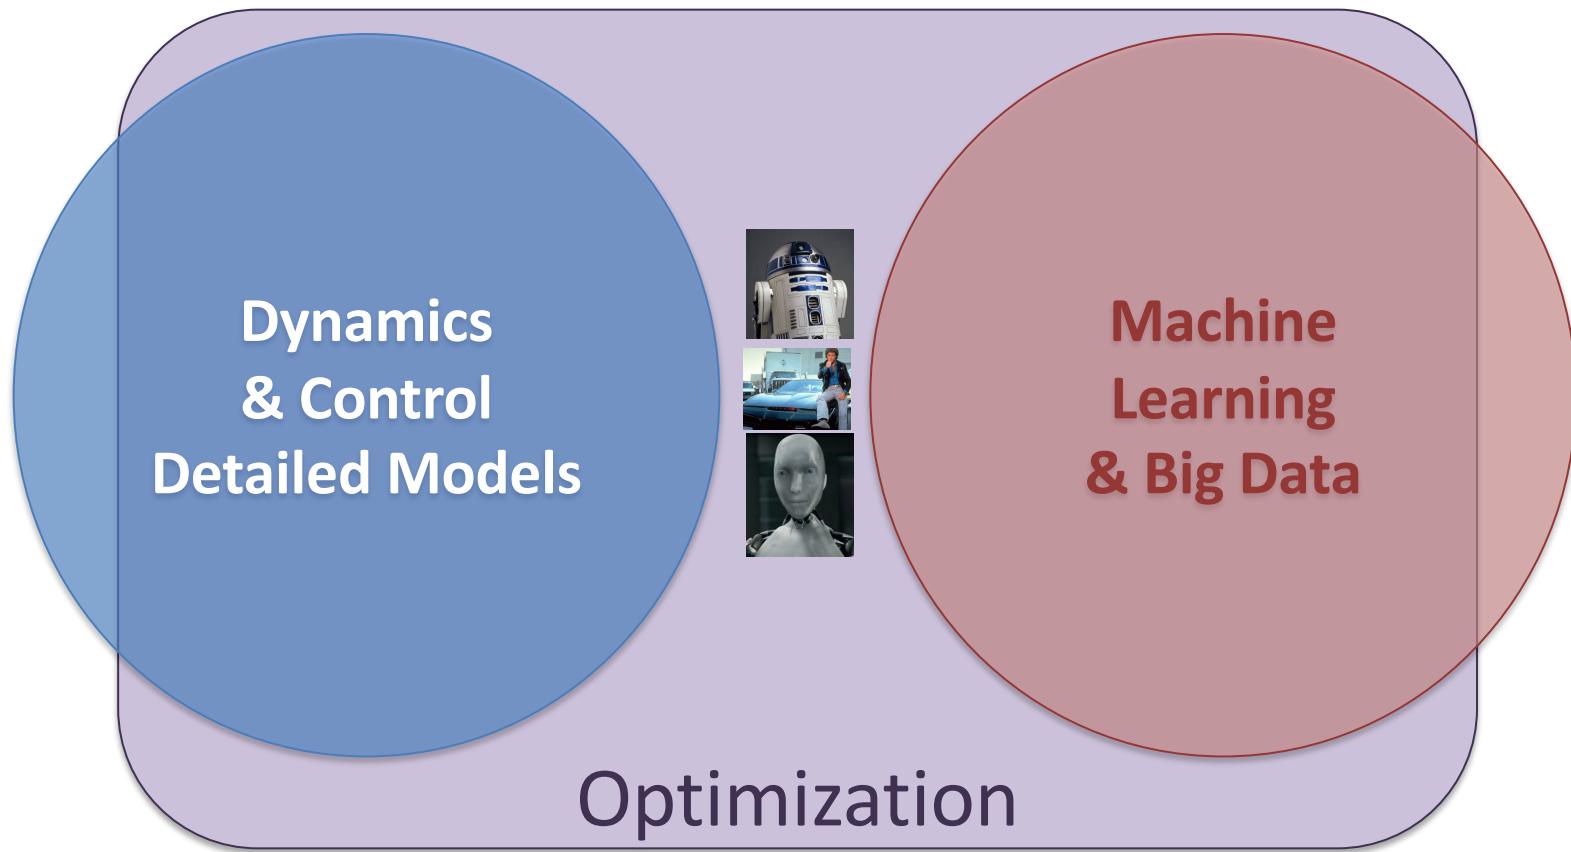


Uber running a red

How do we fix things?



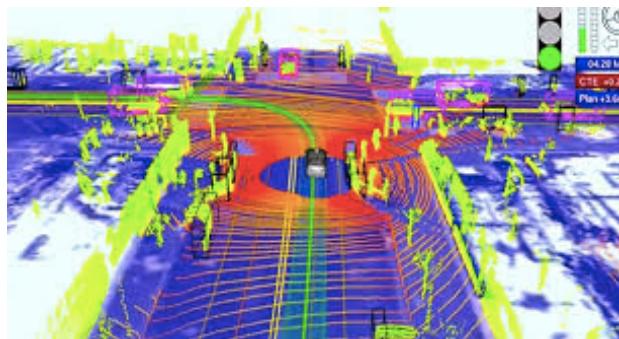
How do we fix things?



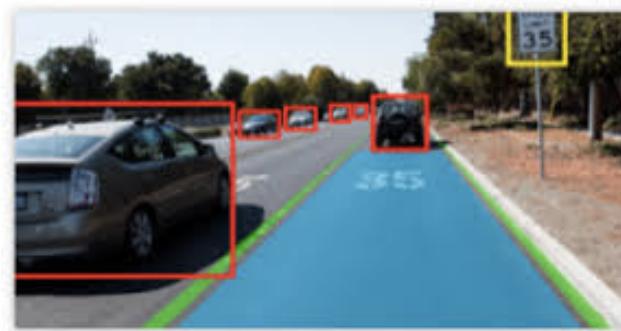
What is ML? Why do we need it?

using past **data** to **learn** about and/or **act** upon the world

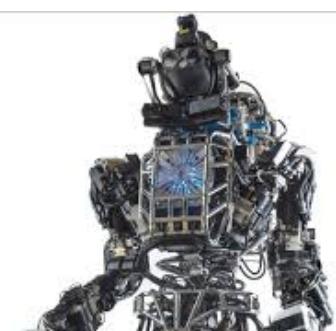
Environments
too complex



Sensing
too complex



Models
too complex

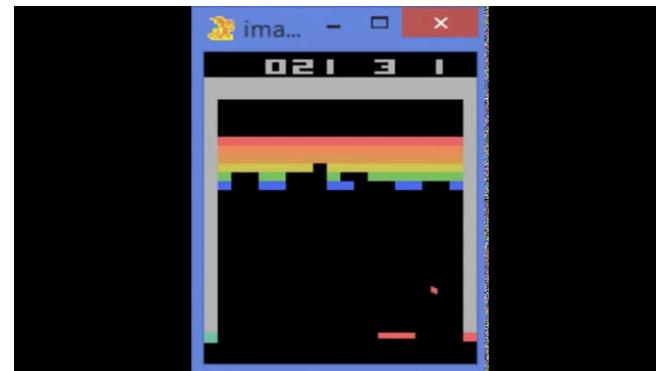


Guarantees are (usually) **probabilistic**

We don't need models!



DeepMind AlphaGo



DeepMind DQN

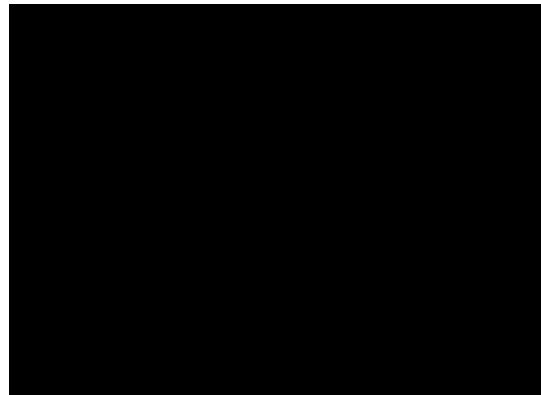


DeepMind AlphaZero

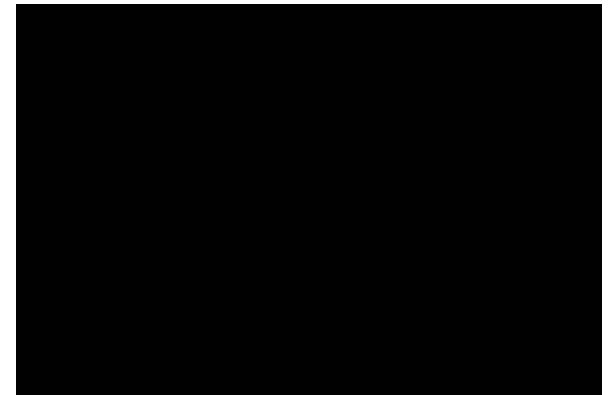
We don't need models!



Goldberg lab, UC Berkeley



Google DeepMind



Open AI

Controlled settings with no cost to failure

Uber's Self-Driving Cars Were Struggling Before Arizona Crash

Tesla Says Autopilot Was Engaged in Fatal Crash Under Investigation in California

Vehicle's system shows driver had hands off the wheel for six seconds before striking highway divider

Las Vegas' self-driving bus crashes in first hour of service

Google AI looks at rifles and sees helicopters

Street sign hack fools self-driving cars

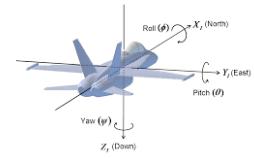
ML deployed in the wild needs guarantees of stability, performance, robustness, safety

Robustness through feedback

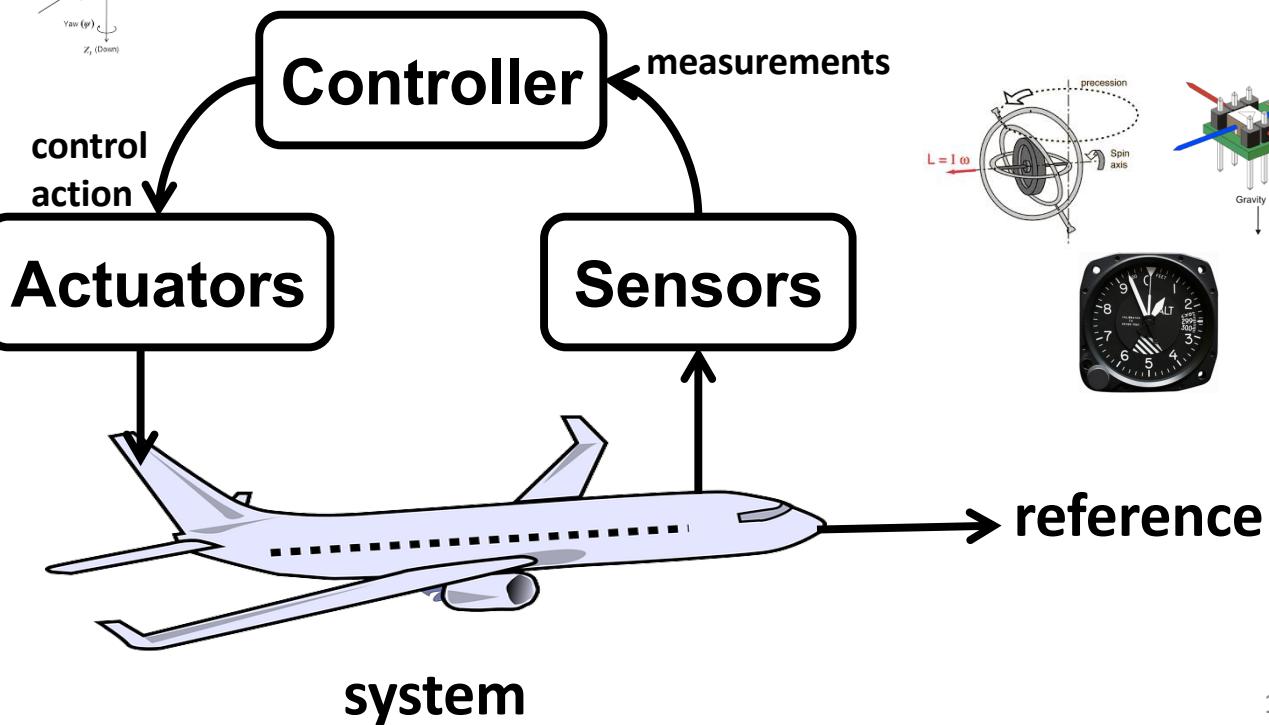
$$\begin{aligned}
 m \left[\frac{dv}{dt} + qv \cdot rv \right] &= F_x - mg \sin \Theta \\
 m \left[\frac{dv}{dt} + ru - pv \right] &= F_y + mg \cos \Theta \sin \Phi \\
 m \left[\frac{dw}{dt} + pw - qu \right] &= F_z + mg \cos \Theta \cos \Phi \\
 M_x &= I_{xx} \frac{dp}{dt} - I_{xz} \frac{dr}{dt} + [I_{xx} - I_{yy}] qr - I_{xz} pq \\
 M_y &= I_{yy} \frac{dq}{dt} + [I_{xx} - I_{zz}] rp + I_{xz} [p^2 - r^2] \\
 M_z &= -I_{xz} \frac{dp}{dt} + I_{zz} \frac{dr}{dt} + [I_{yy} - I_{xx}] pq + I_{xz} qr \\
 p &= \frac{d\Phi}{dt} - \frac{d\Psi}{dt} \sin \Theta \\
 q &= \frac{d\Theta \cos \Phi}{dt} + \frac{d\Psi}{dt} \sin \Phi \cos \Theta \\
 r &= \frac{d\Theta \sin \Phi}{dt} + \frac{d\Psi}{dt} \cos \Phi \cos \Theta
 \end{aligned}$$



disturbances



feedback to close the loop



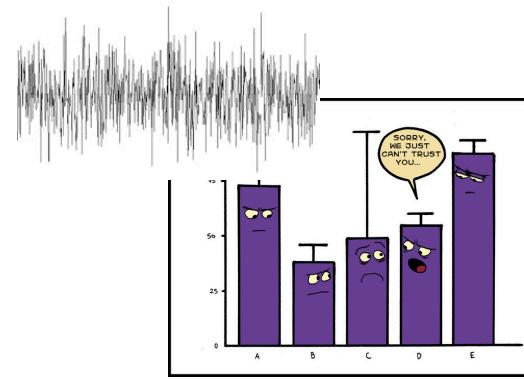
What is Robust Control? Why do we need it?

using **feedback** to **mitigate** the effects of **dynamic uncertainty**

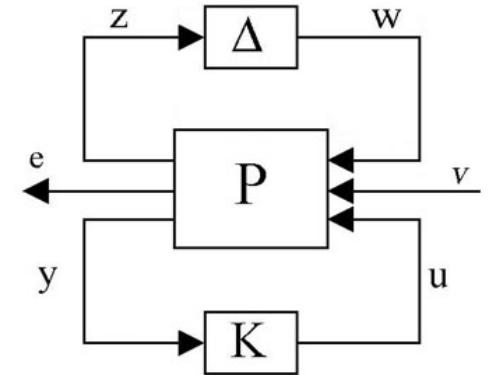
Environments
are uncertain



Sensing/components
are uncertain



Models
are uncertain



Guarantees are (usually) **worst-case and deterministic**

Robustness and Learning?

Machine Learning

uses data to
reduce uncertainty

more data
→ better models/predictions
probabilistic guarantees

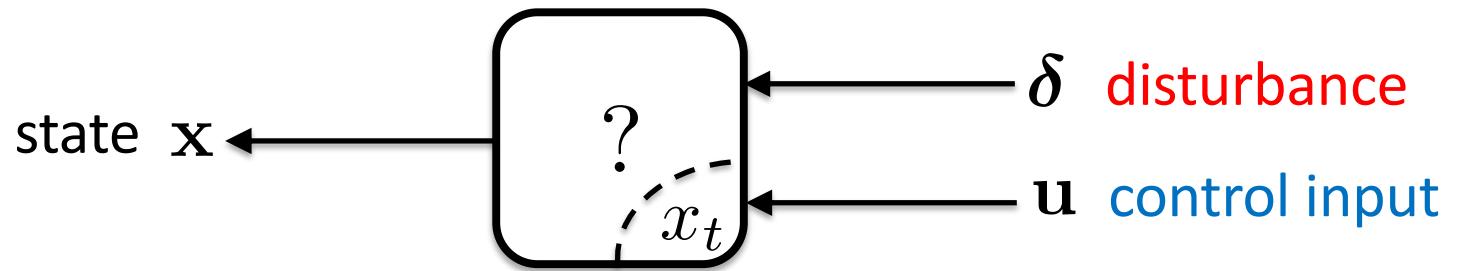
Robust Control

uses feedback to
mitigate uncertainty

better models/predictions
→ better performance
worst-case guarantees

Can ML and RC be combined so that we **safely achieve**
more data → better performance?

Optimal control of an unknown system



Goal:

$$\text{minimize}_{\underline{\pi}_t} \quad \text{Cost}_\delta(x, u) \quad \text{unknown}$$

$$\text{s.t} \quad x_{t+1} = f_t(x_t, u_t, \delta_t)$$

$$\text{state-feedback} \quad u_t = \pi_t(\underline{x_{0:t}}, \underline{u_{0:t-1}})$$

Model Based RL & Control

Uncertainty inherent to the output of a ML algorithm

How do we (safely) learn a model and quantify
its uncertainty?

System Identification, Concentration Bounds, Bootstrapping

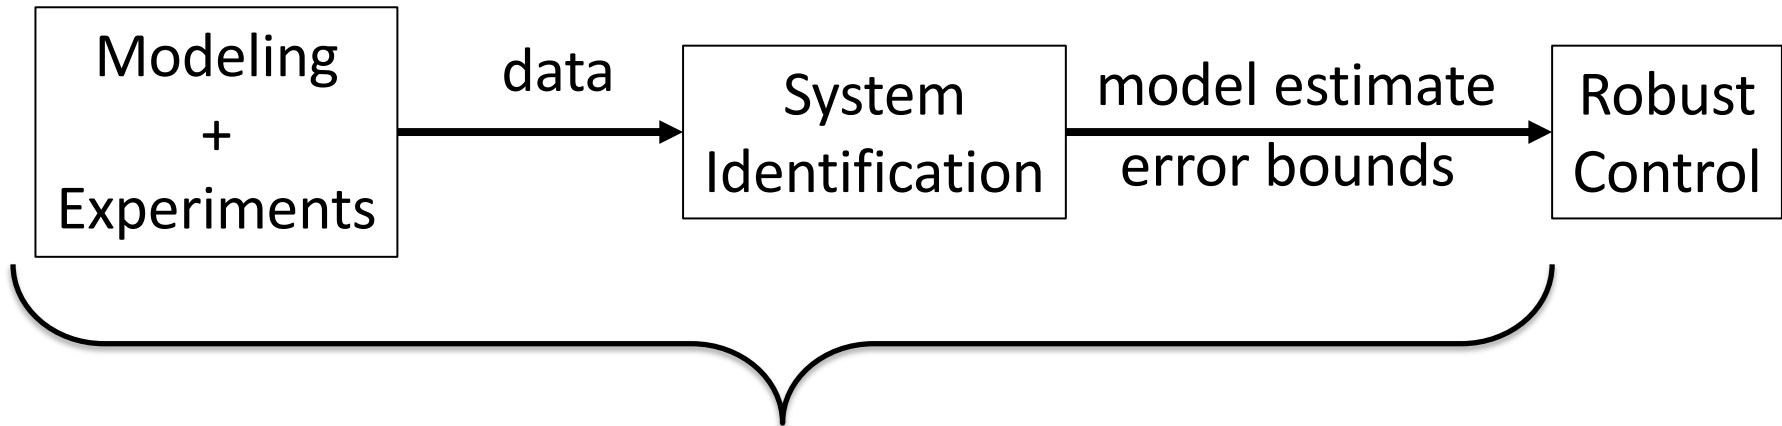
How do we explicitly account for
this uncertainty when we design control policies?

Robust and optimal control, end-to-end analysis

A learning and control pipeline

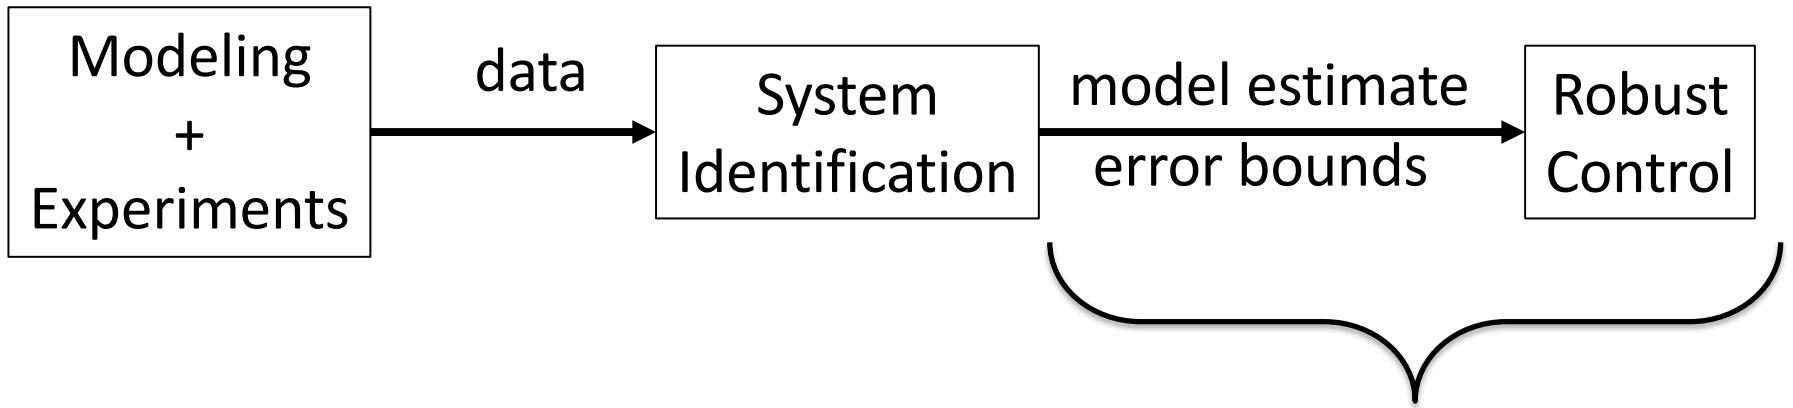


A learning and control pipeline



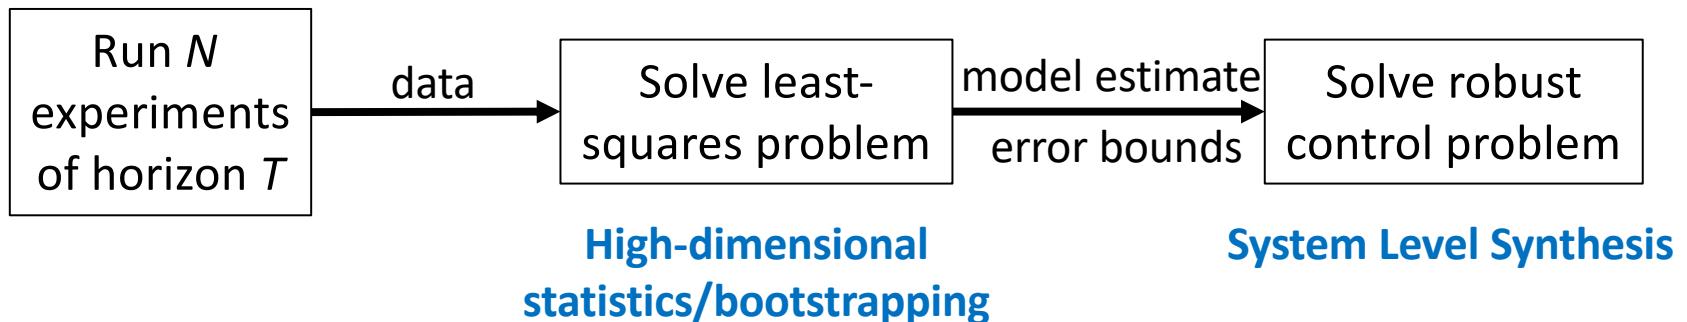
ML & Uncertainty Quantification
show error bounds hold *with high probability*

A learning and control pipeline



RC to Handle Uncertainty
assume error bounds
hold *but are adversarial*

End-to-end guarantees



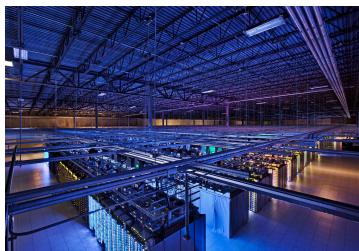
Theorem

With probability $1 - \delta$, for N sufficiently large, the synthesized controller is stabilizing and achieves the relative performance bound

$$\frac{\hat{J} - J_\star}{J_\star} \leq \mathcal{C} \text{ (robustness, excitability)} \sqrt{\frac{(d + p) \log(1/\delta)}{N}}$$

of states N # of inputs

Example: Cooling a server farm



$$\underline{x}_{t+1} = \begin{bmatrix} 1.01 & 0.01 & 0 \\ 0.01 & 1.01 & 0.01 \\ 0 & 0.01 & 1.01 \end{bmatrix} \underline{x}_t + \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \underline{u}_t + \underline{\delta}_t$$

temperature deviations

cooling

“noise”

The equation shows the state transition x_{t+1} as a function of the current state x_t , a control input u_t , and noise δ_t . Red annotations point to the first matrix as "temperature deviations", the second matrix as "cooling", and the third matrix as "noise".

Example: Cooling a server farm



$$x_{t+1} = \begin{bmatrix} 1.01 & 0.01 & 0 \\ 0.01 & 1.01 & 0.01 \\ 0 & 0.01 & 1.01 \end{bmatrix} x_t + \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} u_t + \delta_t$$

Goal: $\min_{u_0, u_1, \dots} \frac{1}{T} \sum_{t=0}^T \mathbb{E} [\underbrace{x_t^T Q x_t}_{\text{don't burn servers}} + \underbrace{u_t^T R u_t}_{\text{don't burn $$$}}]$

Example: Cooling a server farm

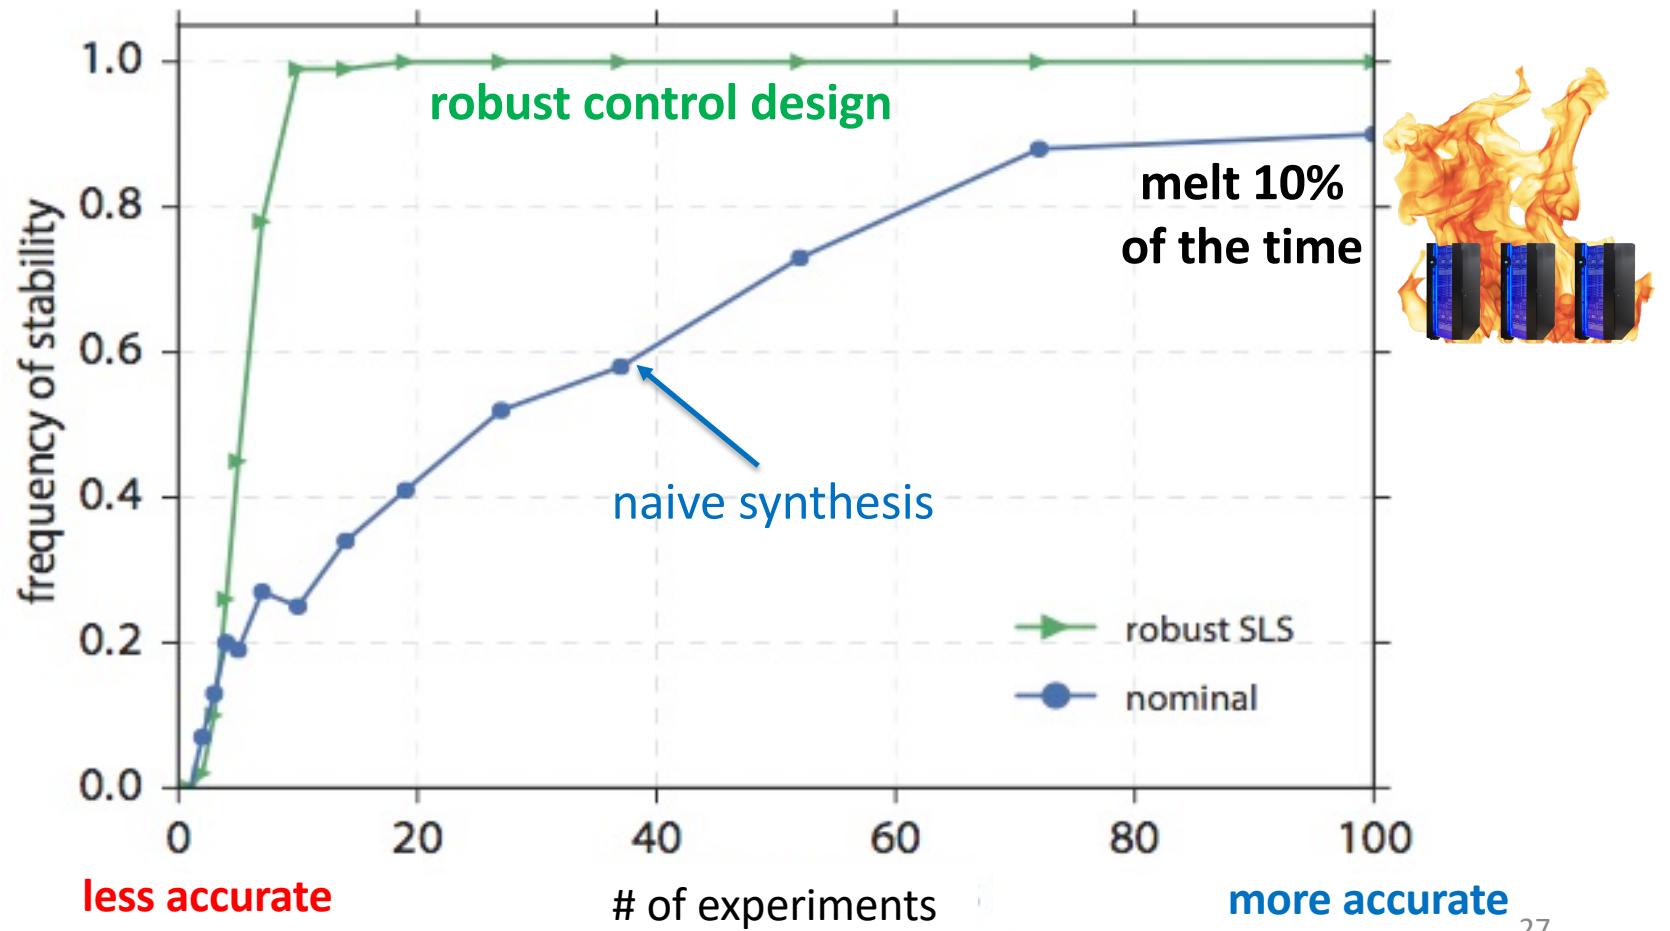


$$x_{t+1} = \begin{bmatrix} ? \end{bmatrix} x_t + \begin{bmatrix} ? \end{bmatrix} u_t + \delta_t$$

Known dynamics: $u_t = K_t x_t$ feedback based policy

Unknown dynamics?

Frequency of servers NOT melting (stability)



Example: Vision based control



Waymo

**How should we train vision systems
for safety-critical control loops?**

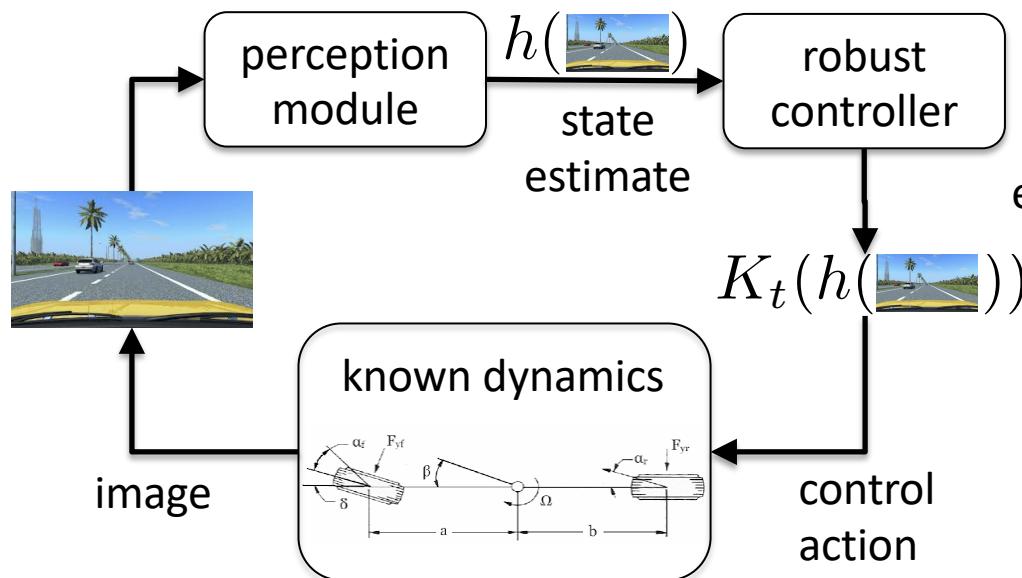
**Should we learn a map from pixels to:
control action (end-to-end)?
waypoints? position?**

Challenges
robustness to environmental shifts
robustness to spurious errors

Vision based control

$$x_{t+1} = Ax_t + Bu_t + w_t$$

$$u_t = K_t(h(\text{image}))$$



ML and Uncertainty Quantification:
learn map h :image → position
and learn **estimation error**

Robust control:
Control robust to estimation error

Challenges
environmental shifts and spurious errors

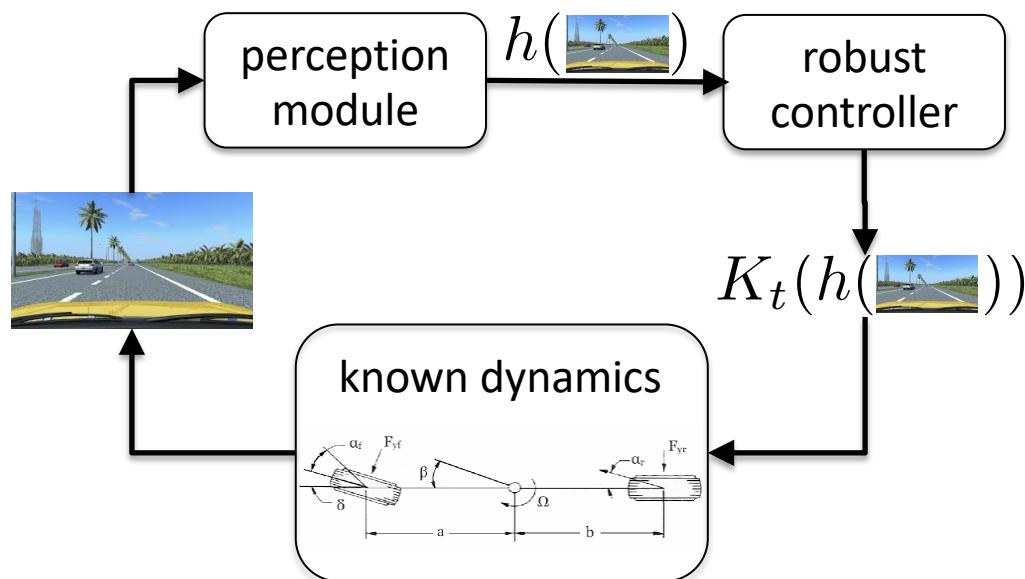
Approach:
take an adversarial view of the world

Can we exploit that we know:
dynamics? control objectives?

Vision based control

$$x_{t+1} = Ax_t + Bu_t + w_t$$

$$u_t = K_t(h(\text{image}))$$



Theorem (informal)

<https://arxiv.org/abs/1907.03680>

Synthesize a **robust controller** with respect to the **learned error model**.

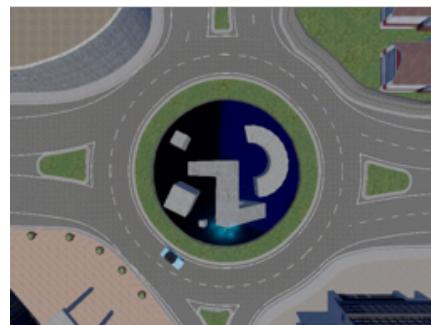
Then* there is a **robust invariant set** around training data for which **generalization error** is bounded, and controller has provable performance.

*under smoothness assumptions

Example: Vision based control

CARLA self-driving
car simulator

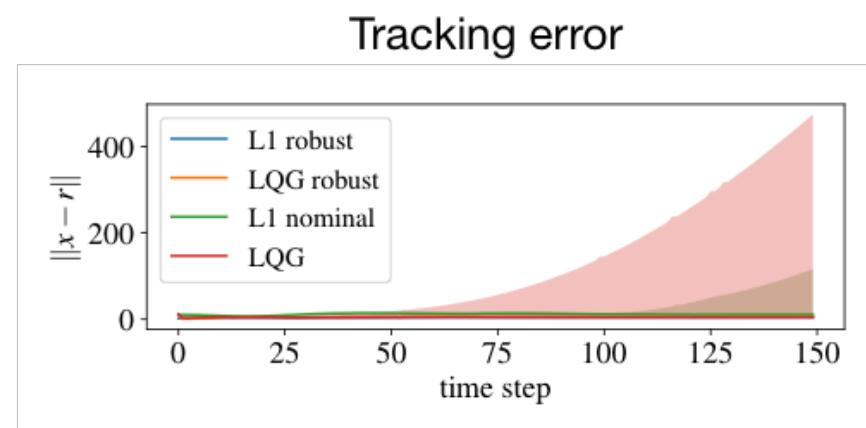
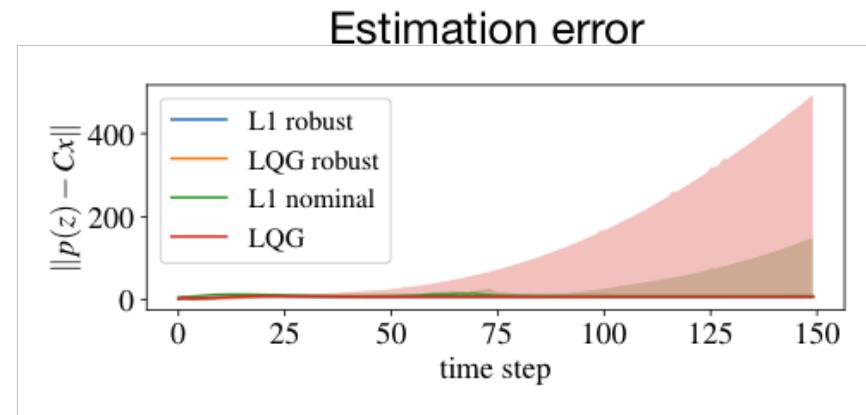
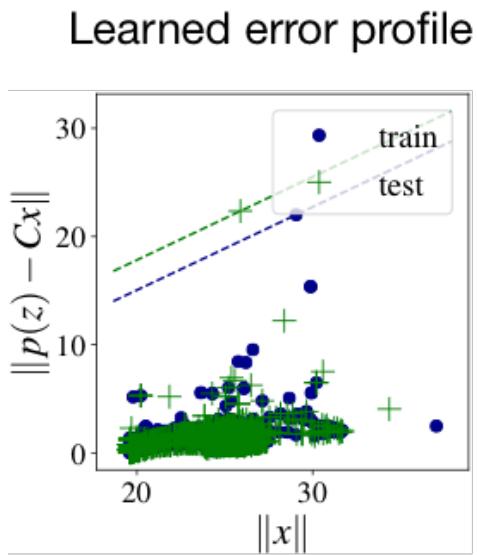
Desired trajectory



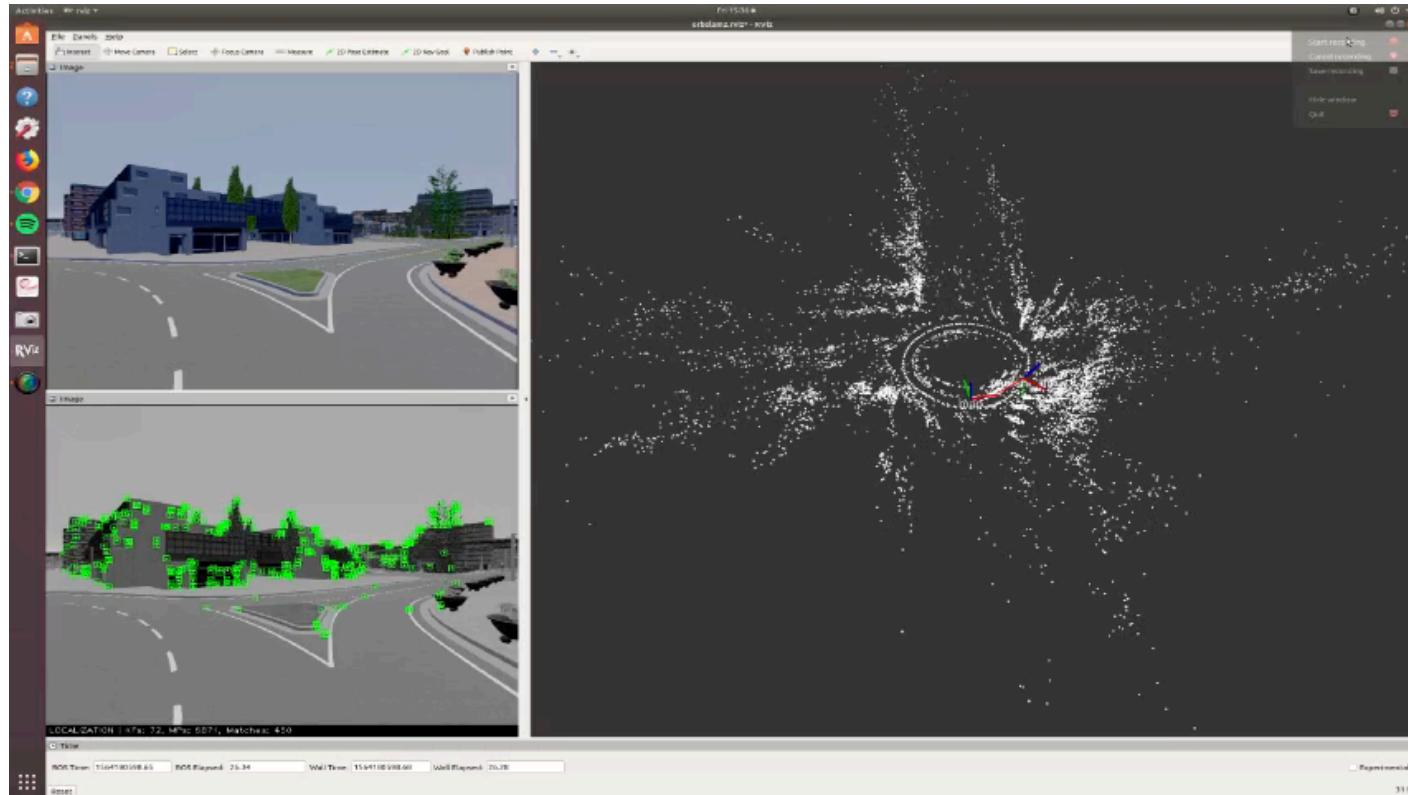
Visual input



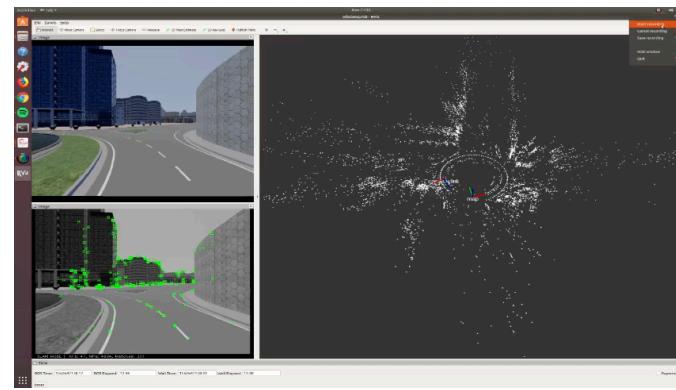
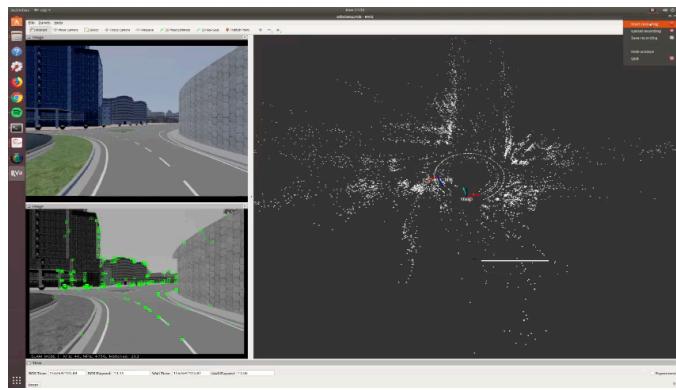
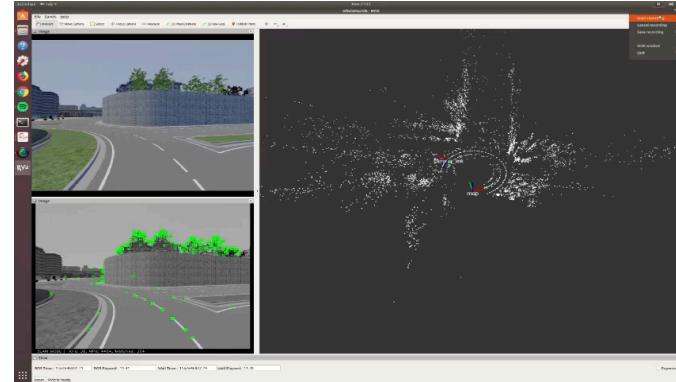
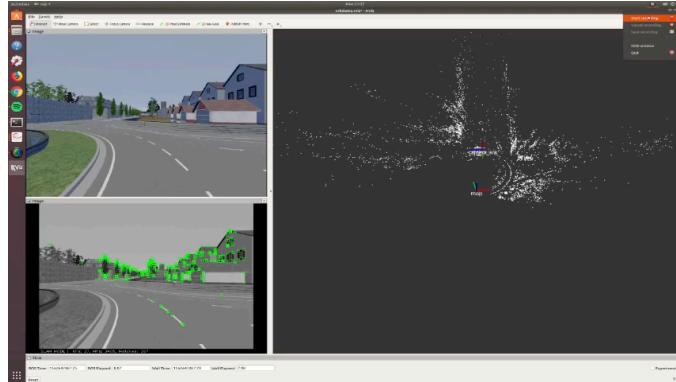
Self-driving car simulator



Self-driving car simulator



Self-driving car simulator



Model Based RL & Control

Uncertainty inherent to the output of a ML algorithm

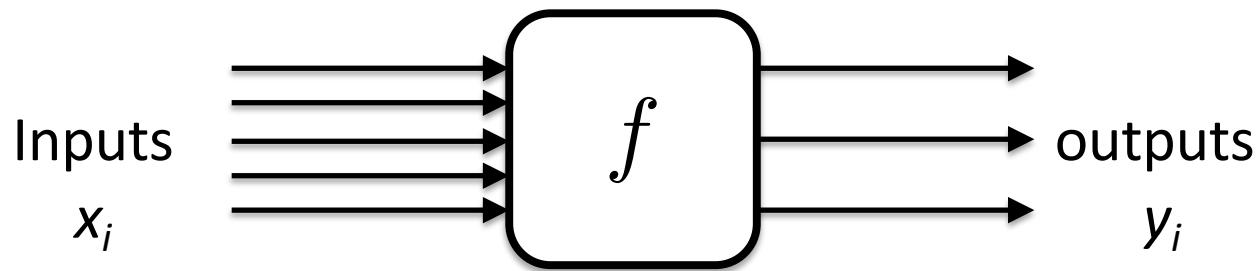
How do we (safely) learn a model and quantify
its uncertainty?

System identification, concentration bounds, bootstrapping

How do we explicitly account for
this uncertainty when we design control policies?

Robust and optimal control, end-to-end analysis

Statistical Learning Theory



Given n training examples: $\{(x_1, y_1), \dots, (x_n, y_n)\} \stackrel{\text{iid}}{\sim} p^n$

Learn function \hat{f} such that

$$\hat{f}(x) \approx y$$

on future (x, y) pairs. We are learning to ***predict*** y from x .

Want to learn a map that ***generalizes*** to future instances.

Example: Linear regression

$$y = \Theta^\top x + w, w \perp x$$

Input x	Output y
High school GPA, SAT score	University GPA
Stock price for FB, NFLX, APPL	NASDAQ idx fund price
Drug dose	Disease population
Current state & input (x_t, u_t)	Next state x_{t+1}

Example: Linear regression

$$y = \Theta^\top x + w, w \sim \mathcal{N}(0, \sigma^2) \perp x$$

Given *iid* n training examples: $\{(x_1, y_1), \dots, (x_n, y_n)\}$

Solve OLS: $\hat{\Theta} = \arg \min_{\Theta} \frac{1}{n} \sum_{i=1}^n (y_i - \Theta^\top x_i)^2$

(assume invertibility) $= \left(\sum_{i=1}^n x_i x_i^\top \right)^{-1} \sum_{i=1}^n x_i y_i$

(plug in $y_i = x_i^\top \Theta + w_i$) $= \Theta + \left(\sum_{i=1}^n x_i x_i^\top \right)^{-1} \sum_{i=1}^n x_i w_i$

estimation error

For *iid* training examples:

- Zero mean (unbiased)
- Converges to 0 (asymptotically consistent)
- Converges whp as $O(n^{-1/2})$ (finite-data guarantees)

Example: Linear System Identification

$$x_{t+1} = Ax_t + Bu_t + w_t, \quad w_t \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma_w^2) \perp x_t$$

Idea: run 1 identification experiment over $n+1$ timesteps and collect n samples $\{(x_t, u_t), x_{\{t+1\}}\}_{t=0}^n$ and solve:

$$(\hat{A}, \hat{B}) = \arg \min_{(A, B)} \sum_{t=0}^n \|x_{t+1} - Ax_t - Bu_t\|_2^2$$

Same problem as last slide, but now data is *not iid*!
This can cause problems: this estimator is not always consistent!

A general framework: Risk Minimization

Want to solve: $\min_{f: \mathcal{X} \rightarrow \mathcal{Y}} \mathbb{E}_p [\ell(y, f(x))]$ (RM)

- $\ell: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ is the loss function (e.g. logistic, least-squares, etc.)
- P is a probability measure on $\mathcal{X} \times \mathcal{Y}$, and is *unknown*
- Only given $S = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\} \sim P^n$
(n.b., this iid assumption will not always hold for us)
- What makes this hard to solve?
 - Only information we have about P is from samples S
 - ∞ -dimensional optimization problem over function space

Empirical Risk Minimization

Want to solve: $\min_{f: \mathcal{X} \rightarrow \mathcal{Y}} \mathbb{E}_p [\ell(y, f(x))]$ (RM)

Instead solve: $\min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i))$ (ERM)

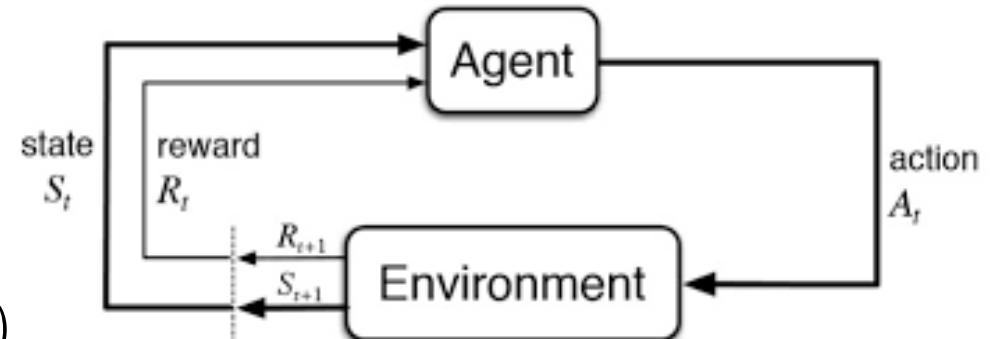
Much of statistical learning theory:
understanding when (ERM) is a good proxy for (RM),
i.e., when small empirical risk implies small population risk

Model free RL & control

$$\text{maximize } \mathbb{E}_{\delta_t} \left[\sum_{t=0}^N R(x_t, u_t) \right]$$

$$\text{subject to } x_{t+1} = f(x_t, u_t, \delta_t)$$

$$u_t = \pi_t(u_{0:t-1}, x_{0:t})$$



Dynamic programming: if we know f , can solve* via DP

$$\mathcal{Q}(x, u) := \max \left\{ \mathbb{E}_{\delta_t} \left[\sum_{t=0}^N R(x_t, u_t) \right] : x_{t+1} = f(x_t, u_t, \delta_t), (x_0, u_0) = (x, u) \right\}$$

Recursively solve Bellman equation:

$$\mathcal{Q}_N(x, u) = R(x, u), \quad \mathcal{Q}_k(x, u) = R(x, u) + \mathbb{E}_{\delta} \left[\max_{u'} \mathcal{Q}_{k+1}(f(x, u, \delta), u') \right]$$

Optimal control action: $u_k = \arg \max_u \mathcal{Q}_k(x_k, u)$

Reinforcement vs Machine Learning

Machine learning

Draw n training examples: $\{(x_1, y_1), \dots, (x_n, y_n)\} \stackrel{\text{iid}}{\sim} p^n$
Learn function f such that $f(x) = \hat{y} \approx y$

Reinforcement learning

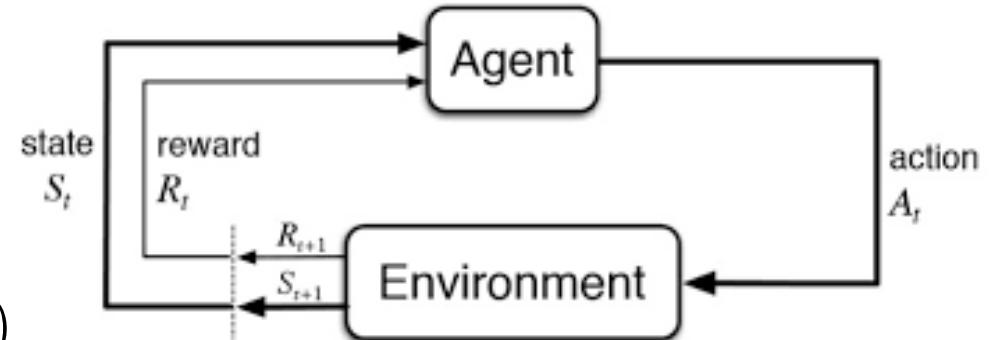
Generate sample trajectories: $\{(u_{0:t-1}^{(i)}, x_{0:t}^{(i)}, r_{0:t}^{(i)})\} \stackrel{\text{iid}}{\sim} p_{\delta, \pi_{train}}^n$
Learn policy π that produces u so that r is large
Policy can explore to learn about system, or use existing knowledge
to control: this leads to *exploration vs. exploitation tradeoff*

Model free RL & control

$$\text{maximize } \mathbb{E}_{\delta_t} \left[\sum_{t=0}^N R(x_t, u_t) \right]$$

$$\text{subject to } x_{t+1} = f(x_t, u_t, \delta_t)$$

$$u_t = \pi_t(u_{0:t-1}, x_{0:t})$$



Approximate Dynamic Programming: if we don't know f , if state/action space is too big, etc.

Draw a sample trajectory $\{x_k, u_k\}$ by playing a policy – if it is optimal, then we have

$$Q(x, u) \approx R(x, u) + \max_{u'} Q(x_{k+1}, u')$$

(approximately in expectation)

Suggests update rule:

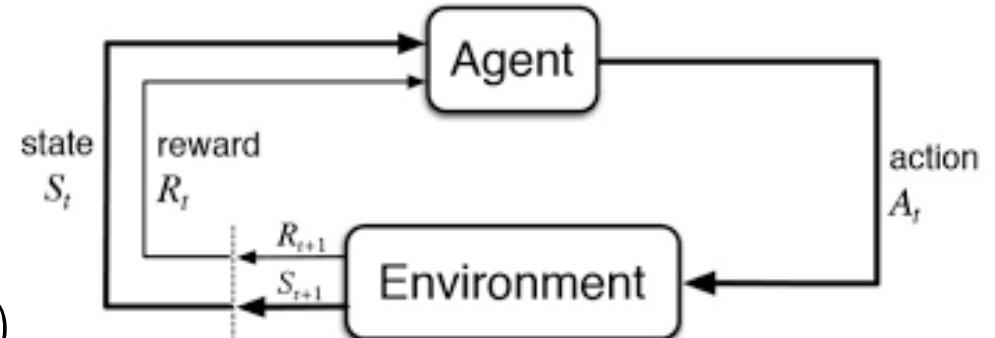
$$Q_{new}(x_k, u_k) = (1 - \eta)Q_{old}(x_k, u_k) + \eta \left(R(x_k, u_k) + \max_{u'} Q_{old}(x_{k+1}, u') \right)$$

Model free RL & control

$$\text{maximize } \mathbb{E}_{\delta_t} \left[\sum_{t=0}^N R(x_t, u_t) \right]$$

$$\text{subject to } x_{t+1} = f(x_t, u_t, \delta_t)$$

$$u_t = \pi_t(u_{0:t-1}, x_{0:t})$$



Direct Policy Search: if we don't know f , if state/action space is too big, etc.,
We can instead optimize directly over policy π

Parametrize probabilistic policy as $\pi(x_t) = p(u|u_{0:t-1}, x_{0:t}, \theta)$ for θ our design parameter
Sample trajectories $(u_{0:t-1}, x_{0:t})$, and then run:

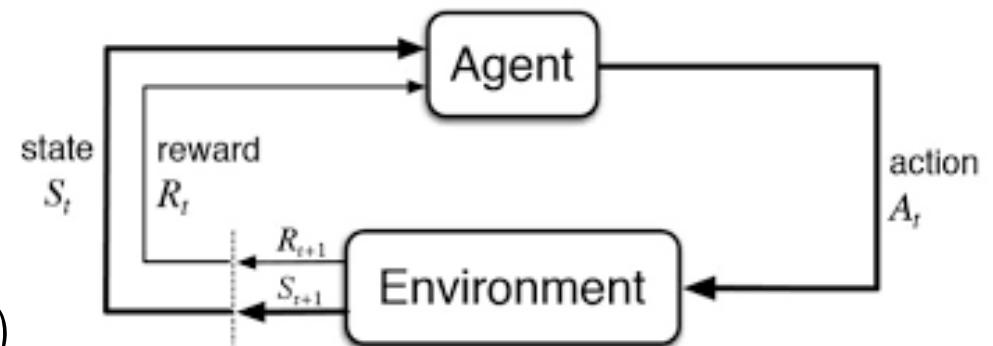
- policy gradient updates using stochastic gradient descent (SGD)
- random search via derivative-free-optimization (DFO)

What about Deep Learning?

$$\text{maximize } \mathbb{E}_{\delta_t} \left[\sum_{t=0}^N R(x_t, u_t) \right]$$

$$\text{subject to } x_{t+1} = f(x_t, u_t, \delta_t)$$

$$u_t = \pi_t(u_{0:t-1}, x_{0:t})$$



Model based: parameterize learned model by a neural net

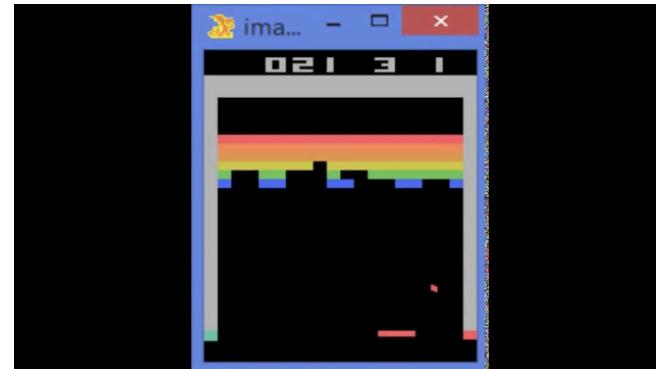
ADP: parameterize Q-function by a neural net

Direct policy search: parameterize policy distribution by a neural net

We don't need models!



DeepMind AlphaGo

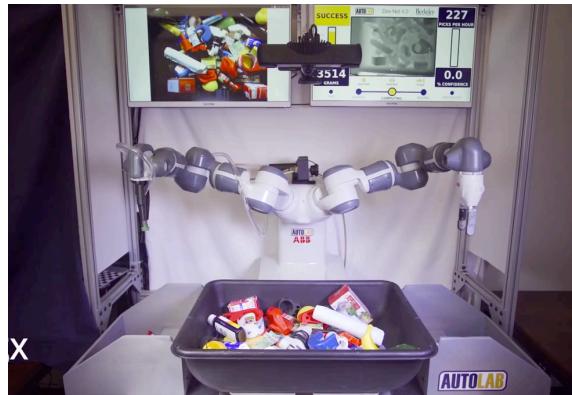


DeepMind DQN

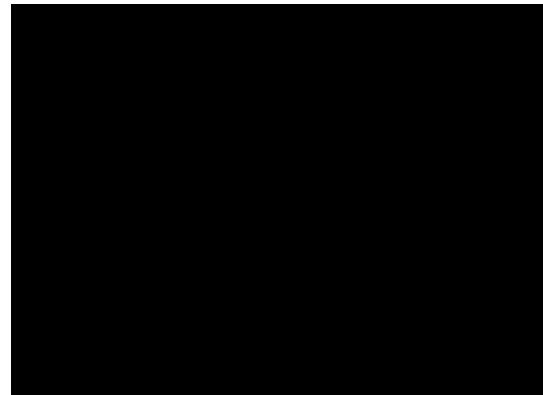


DeepMind AlphaZero

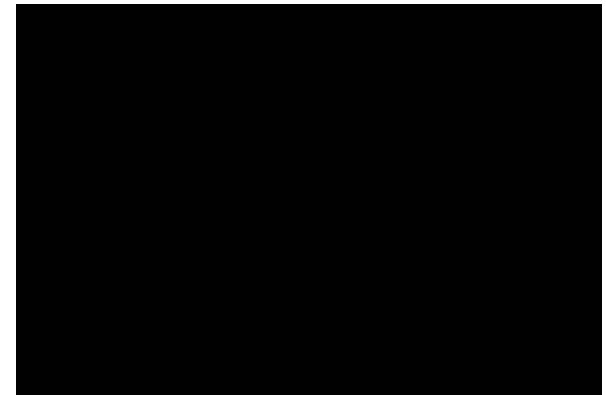
We don't need models!



Goldberg lab, UC Berkeley



Google DeepMind



Open AI

Controlled settings with no cost to failure

Model Free Control

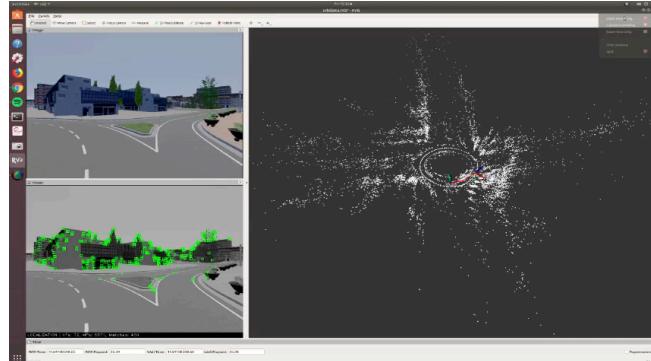
What is the right function space to work with?

Can we enforce stability in a model-free setting?
Safety? Robustness?

How much data do we need to learn good control policies?

When should we use model-free vs. model-based methods?

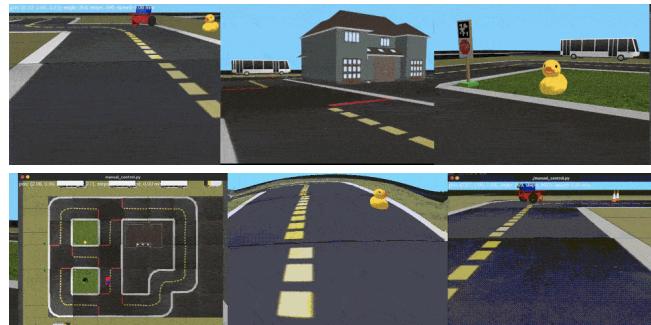
Simulation environments



CARLA



FLIGHTGOOGLES



DUCKIETOWN

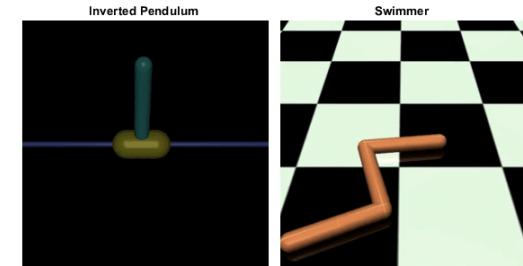


UDACITY/UNITY-3D

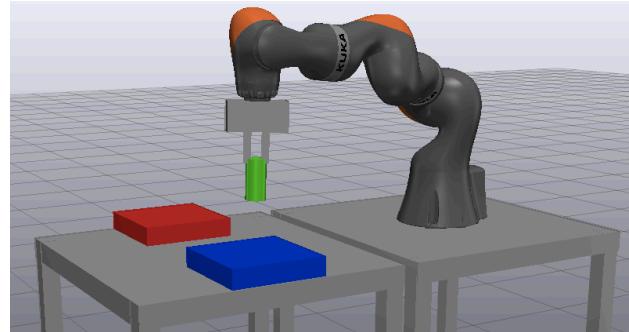
Simulation environments



MuJoCo



OpenAI Gym



Drake

Recap

- **Learning** to deal with complexity, **control** to deal with uncertainty
- **Learning + control** promises high-performing and adaptive systems, **but** stability/safety/robustness must be established!
- **Model based approaches**
Learn model and quantify uncertainty, then apply robust control: can we get end-to-end performance bounds?
- **Model free approaches**
RL/ADP/PG: broadly applicable, easy to use, impressive demos, but what about safety? robustness? sample efficiency?

Please do the following!

- Please fill out background poll at <https://www.surveymonkey.com/r/XFSR26N>
- **If you are sitting in, e-mail me with subject line starting with ESE680, and please still fill this out.**
- Sign up sheet: https://docs.google.com/spreadsheets/d/1Kqa-ZnnqO-r7-c0rTKNSaCNiz_hTQARH9V_Ug8blp3Q/edit?usp=sharing
- Need someone to scribe next class, and please sign up for presentations on Sep 10 and 12 and e-mail me to set up a meeting time next week!